

# Classical text mining

5/5 points (100.00%)

Quiz, 5 questions

✓ **Congratulations! You passed!**

Next Item



1 / 1  
points

1.

---

Choose true statements about text tokens.



Lemmatization is always better than stemming



**Un-selected is correct**



Stemming can be done with heuristic rules



**Correct**

Yeah, Porter stemmer works this way.



Lemmatization needs more storage than stemming to work



**Correct**

This is true, you have to store information about all possible word forms in the vocabulary.



A model without stemming/lemmatization can be the best



**Correct**

This is true. Word2vec embeddings, for instance, are trained on raw tokens.

## Classical text mining

5/5 points (100.00%)

Quiz, 5 questions



1 / 1  
points

2.

Imagine you have a texts database. Here are stemming and lemmatization results for some of the **words**:

Word	Stem	Lemma
<b>operate</b>	oper	operate
<b>operating</b>	oper	operating
<b>operates</b>	oper	operates
<b>operation</b>	oper	operation
<b>operative</b>	oper	operative
<b>operatives</b>	oper	operative
<b>operational</b>	oper	operational

Imagine you want to find results in your texts database using the following queries:

1. **operating system** (we are looking for articles about OS like Windows or Linux)
2. **operates in winter** (we are looking for machines that can be operated in winter)

Before execution of our search we apply either stemming or lemmatization to both query and texts. Compare stemming and lemmatization for a given query and choose the correct statements.



Stemming provides higher recall for **operates in winter** query.



**Correct**

This is true, lemmatization would only find exact matches with **operates** and lose a lot of relevant forms like **operational**.



Stemming provides higher F1-score for **operating system** query.

## Classical text mining

5/5 points (100.00%)

**Un-selected is correct**

Quiz, 5 questions



Stemming provides higher precision for **operating system** query.



**Un-selected is correct**



Lemmatization provides higher precision for **operates in winter** query.



**Correct**

This is true, but it would loose a lot of other relevant forms.



1 / 1  
points

---

**3.**

Choose correct statements about bag-of-words (or n-grams) features.



Hashing **vectorizer** (object that does vectorization) needs an amount of RAM proportional to vocabulary size to operate.



**Un-selected is correct**



Classical bag-of-words **vectorizer** (object that does vectorization) needs an amount of RAM at least proportional to  $T$ , which is the number of unique tokens in the dataset.



**Correct**

This is true, you have to store a hash map {token: index} to be able to vectorize new texts.

## Classical text mining

5/5 points (100.00%)

☐ For bag-of-words features you need an amount of RAM at least proportional to  $N \times T$ , where  $N$  is the number of documents,  $T$  is the number of unique tokens in the dataset.



**Un-selected is correct**



We prefer **sparse** storage formats for bag-of-words features.



**Correct**

This is true. We have a lot of zeros in these features, that's why we can store them efficiently in sparse formats (look at `sklearn.feature_extraction.text.TfidfVectorizer` and `scipy.sparse.csr.csr_matrix`).



You get the same vectorization result for any words permutation in your text.



**Un-selected is correct**



1 / 1  
points

4.

Let's consider the following texts:

- good movie
- not a good movie
- did not like
- i like it

# Classical text mining

5/5 points (100.00%)

Quiz, 5 questions

Let's count **Term Frequency** here as a distribution over tokens in a particular text, for example for text "good one" we have  $TF = 0.5$  for "good" and "one" tokens.

## Term frequency (TF)

- $tf(t, d)$  – frequency for term (or n-gram)  $t$  in document  $d$
- Variants:

weighting scheme	TF weight
binary	0, 1
raw count	$f_{t,d}$
term frequency	$f_{t,d} / \sum_{t' \in d} f_{t',d}$
log normalization	$1 + \log(f_{t,d})$

## Inverse document frequency (IDF)

- $N = |D|$  – total number of documents in corpus
- $|\{d \in D: t \in d\}|$  – number of documents where the term  $t$  appears
- $idf(t, D) = \log \frac{N}{|\{d \in D: t \in d\}|}$

What is the **sum** of TF-IDF values for 1-grams in "good movie" text? Enter a math expression as an answer. Here's an example of a valid expression:  $\log(1/2)*0.1$ .

Preview

$$-0.5 \log(3) - 0.5 \log(2) + 1.0 \log(5)$$

# Classical text mining

**5/5 points (100.00%)**

Quiz, 5 questions

Your answer,  $0.5 \cdot \log(5/3) + 0.5 \cdot \log(5/2)$ , is equivalent to the instructor's answer ( $0.5 \cdot \log(5/3) + (0.5 \cdot \log(5/2))$ ).



1 / 1  
points

**5.**

What models are usable on top of bag-of-words features (for 100000 words)?



Naive Bayes

**Correct**

SVM

**Correct**

Decision Tree

**Un-selected is correct**

Gradient Boosted Trees

**Un-selected is correct**



Logistic Regression

# Classical text mining

**5/5 points (100.00%)****Correct**

Quiz, 5 questions

